

# Localized Optimizations over Path Segments

IETF 106, LOOPS side meeting 2019-11-19

# Note Well

- You may be recorded
- Be nice, and be professional
- The IPR guidelines of the IETF apply:  
see <http://ietf.org/ipr> for details.

# Agenda

- What is LOOPS (25)
- Technical discussion about specific items (3x15)
- Next steps (5)

# What is LOOPS?

**Localized Optimizations  
over Path Segments**

IETF 106, LOOPS side meeting 2019-11-19  
(Slides compiled by Carsten Bormann)



Packet loss

Packet loss


Packet loss



Tail  
loss!



# LOOPS Opportunity



# Recover Packets Locally

Reduce end-to-end packet loss

Recover locally, **where needed**, with low latency

In the **network**

Host participation not required



# Don't look

# Don't touch

Works with any kind of IP packets

# How to recover?

- **Retransmission**
  - Reverse information needed: ACK/NACK
  - Forward information: sequence numbering (if needed)
- **Forward Error Correction** (redundancy)
  - Can use dynamic selection of block size/rate: measurement input
  - “Retransmission” also possible by adding FEC
- Aim for low setup overhead
- Keep most setup out of protocol (“controller model”)

← Piggyback (Tunnel), separate Packets

← Tunnel

# How not to blow up the Internet

- Concealing losses removes important congestion signal
  - End-hosts would ramp up to higher rates, increase congestion
- Need **congestion feedback**
  - Preferred: ECN
  - Fallback: Selective dropping (selective recovery, actually)
- Host transport protocol improvements will help improve LOOPS performance, but are not prerequisite to obtaining benefit

# Elements of LOOPS

- Information model for local **recovery**: in-network retransmission/FEC
  - Can be encapsulated in a variety of formats; define some of those
- Local **measurement**: e.g. segment forward delay/variation
  - To set recovery parameters
  - To determine if loss was caused by congestion
- Congestion **feedback**:  
ECN (or drops) to inform end hosts about congestion loss

# Freezer (not currently in scope)

- Multipath
- Measurement along string of LOOPS pairs (“almost e2e”)
- MTU handling, fragmentation, aggregation, header compression
- Selection of one or more specific tunnel encapsulation or measurement formats  
(beyond “sketches” showing it can be made to work)



# LOOPS vs. transport protocols

- LOOPS is separate from the end-to-end transport protocol
  - Hands-off approach: don't meddle
  - Do not assume the end-to-end protocol is out to help us, either
  - No direct control over sending rate (cc feedback only)
- LOOPS should not just be a classical transport protocol
  - Residual loss is OK
  - More choices: Tight interaction with the path segment being optimized

# Where “transport protocol” intuition may not even work

- Relatively controlled/managed environment; setup mechanism assumed (can supply parameters so not everything needs to be high dynamic range)
- No full reliability intended; remaining gaps are OK (and at some point must leave the focus of attention)
  - Setup might set upper bound for overhead volume (e.g., 10 %), can well be “risky” in the way that this is used
- Tunnels usually have packets in flight (possibly a large number); tail processing rarely invoked (but may still be desired); don’t need overly conservative RTO

# Documents out there

- Use cases and problem statement: “LOOPS (Localized Optimizations on Path Segments) Problem Statement and Opportunities for Network-Assisted Performance Enhancement”  
<draft-li-tsvwg-loops-problem-opportunities>
- Protocol: “LOOPS Generic Information Set” <draft-welzl-loops-gen-info>
- One of the Encapsulations: “Embedding LOOPS in Geneve”  
<draft-bormann-loops-geneve-binding-00.txt>
- Charter proposal for a LOOPS WG <<https://github.com/loops-wg/charter>>
- LOOPS mailing list loops@ietf.org

# Related work (see IETF105 BOF)

- Encapsulations: Many (e.g., NVO3 for Geneve; GUE; GRE?)
- FEC: NWCRG for e.g., sliding window FEC, encapsulation techniques
- Tunnel congestion Feedback (TSVWG)
- Also: measurement work, IOAM;  
knowledge about behavior of transport protocols (TCP, QUIC)  
adaptation layer retransmission work (6Lo Fragment Recovery)

# Sliding Window FEC

- Sliding windows fit quite well to LOOPS application  
(Can also use traditional block formats)
- Various drafts for FEC scheme and specific embeddings in NWCRG and TSVWG, e.g.,
  - "Sliding Window Random Linear Code (RLC) Forward Erasure Correction (FEC) Schemes for FECFRAME" <draft-ietf-tsvwg-rlc-fec-scheme-16.txt>
  - "Forward Error Correction (FEC) Framework Extension to Sliding Window Codes" <draft-ietf-tsvwg-fecframe-ext-08.txt>



# (F)AQ (1)

- So this is only about encrypted traffic?
  - Any traffic is welcome, we just don't try to peek beyond L3 info
- So how do you know which packets are worth recovering?
  - Today we don't. If more L3 marking becomes available, we'd use it.
- How do you transport your measurement-related information?
  - Forward info: In encapsulation extension (e.g., with sequence number).  
Reverse info: The same way we transport the ACK channel. Depends on encapsulation.

# (F)AQ (2)

- How do you avoid spending more for LOOPS encapsulation than the performance enhancement is worth?
  - LOOPS will need some management that is weighing this (and doing the pair setup in the first place)
  - Can dynamically switch off and on (e.g., based on monitoring)
- How to relay congestion for non-ECN-capable transports?
  - Dropping. Or, actually, not even requesting a retransmission when a congestion event would be relayed anyway.

**Slides based on**

# **LOOPS Generic Information Set**

**draft-welzl-loops-gen-info-00**

LOOPS BoF

IETF 105 - Montreal

Michael Welzl, U. Oslo

Carsten Bormann, ed., U. Bremen TZI

# Why look at this draft now?

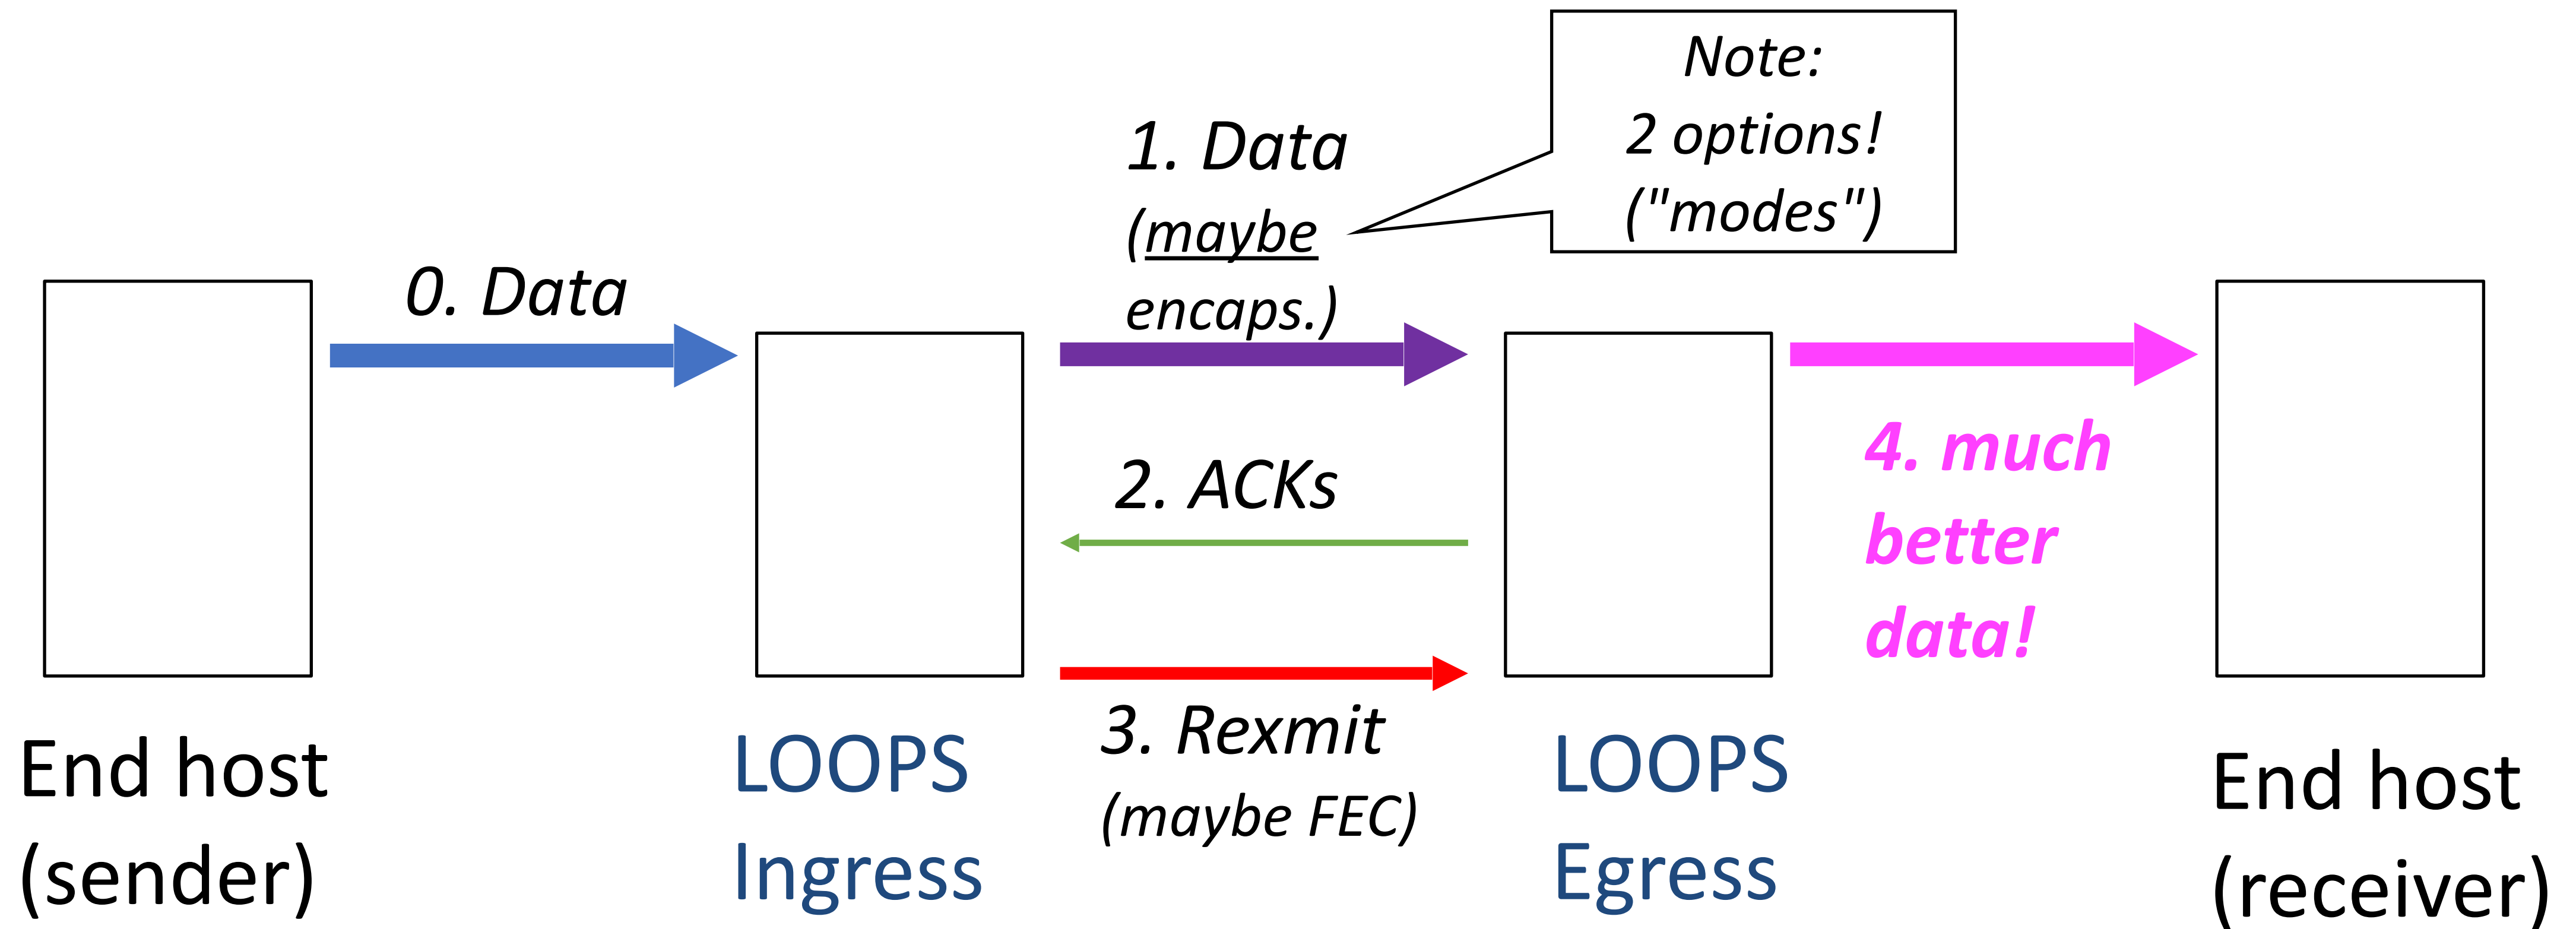
- "The present document is a **strawman** for the set of information that would be interchanged in a LOOPS protocol, without already defining a specific data packet format."

→ an overview of how a LOOPS protocol **could** work

... as an existence proof, and to aid visualization

~~We are not picking alternatives today.~~ We are starting to collect reasons for decisions, to prepare the WG work.

# Context





# Problems to address

- From previous slide:
  - Loss detection/retransmission
  - FEC control
- Also: detect congestion on ingress-egress segment
  - Measurement / congestion detection
  - Congestion signaling to end hosts if congestion was detected
- Next: some concrete ideas on how to deal with these problems
  - Just a strawman (two, actually)!

# Tunnel mode

# 1. Ingress forwards

- **Encapsulate**: not tied to any specific overlay protocol
  - Document contains sketches of bindings to GUE and Geneve
- We don't try to understand data after the IP header
  - Hence, need to add our own **Packet Sequence Number (PSN)**
- Some more information added
  - **Tunnel type**
  - **"ACK desirable" flag** (asks for feedback block 1, next slide)
  - **Anything needed by FEC**

## 2. Egress answers

- **Block 1** (optional, only upon "ACK desirable" — timely feedback)
  - PSN being acknowledged
  - Absolute time of reception for the packet acked (PSN)
- **Block 2** (optional — eventual feedback)
  - an ACK bitmap (based on PSN)
  - a delta indicating the end PSN of the bitmap
- Can be interspersed and repeated
- Can be piggybacked on a reverse direction data packet or sent separately
- Usually aggregated in some useful form

### 3. Ingress retransmits

- Detects need for rexmit via ~~NACK~~ or ~~RTO~~ timing
  - Make decision based on congestion
    - Use ECN if possible
    - Calculate latency variation from timestamps in feedback blocks 1
- ... Or, rather than "just rexmit", send FEC repair packets



## 4. Egress forwards

- De-FEC
- Inform end hosts about congestion if needed
  - Might be able to distinguish "real" congestion from, e.g., corruption loss
  - ECN much preferred as an outgoing signal!

# Summary: information exchanged

- **Forward:** encapsulation, containing...
  - Packet Sequence Number (PSN)
  - Tunnel type
  - "ACK desirable" flag (asks for feedback block 1, next slide)
  - Anything needed by FEC
- **Backward:** optional blocks type 1 (timely) and 2 (eventual)...  
*(can be piggybacked, aggregated, interspersed, repeated, ...)*
  - **Block 1** (optional, only upon "ACK desirable" — timely)
    - PSN being acknowledged
    - Absolute time of reception for the packet acked (PSN)
  - **Block 2** (optional — eventual)
    - an ACK bitmap (based on PSN)
    - a delta indicating the end PSN of the bitmap

# Transparent mode

A bit more radical... describing least intrusive method here:

"Never delay and don't even tunnel"

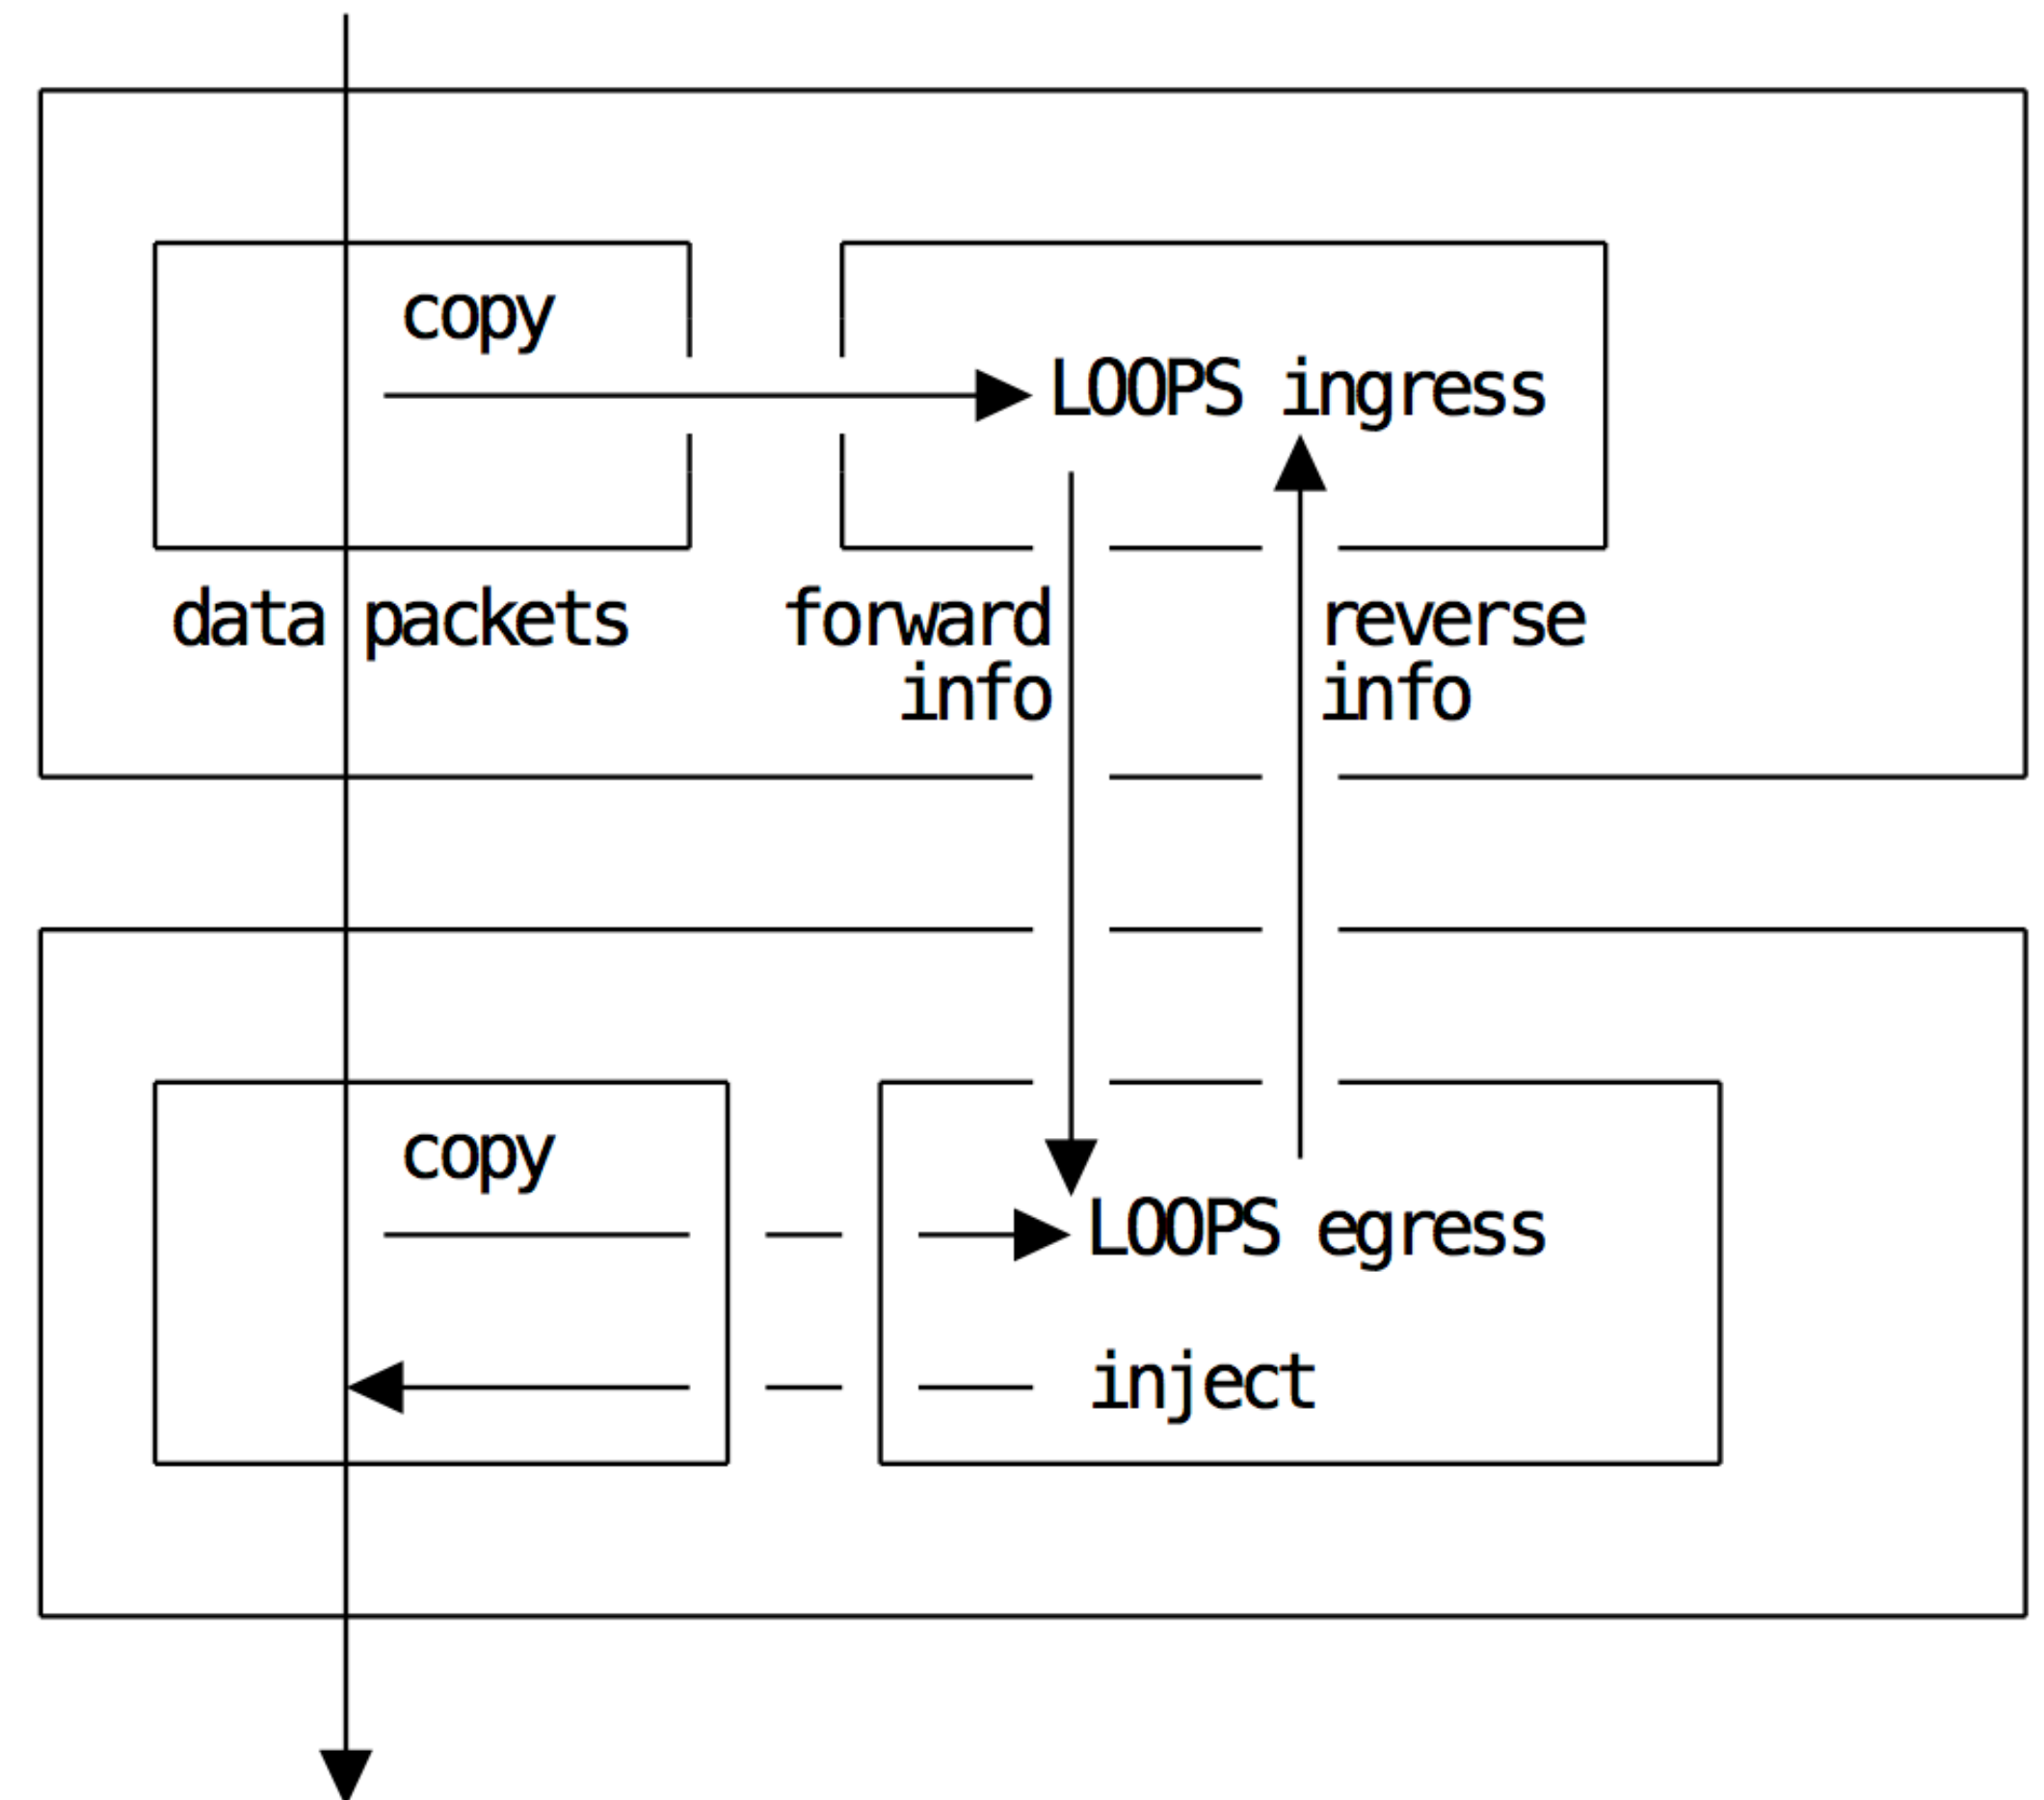
Just discussing rexmit; FEC could also be done

# 1. Ingress forwards

- Just forward
- Also, keep a copy of packets, with a hash for identification
  - From immutable header fields
  - May need to include data beyond IP

# 2. Egress answers

- **ACK everything; no NACK possible**
  - Same hash calculation as ingress
- ACK format similar to tunnel mod

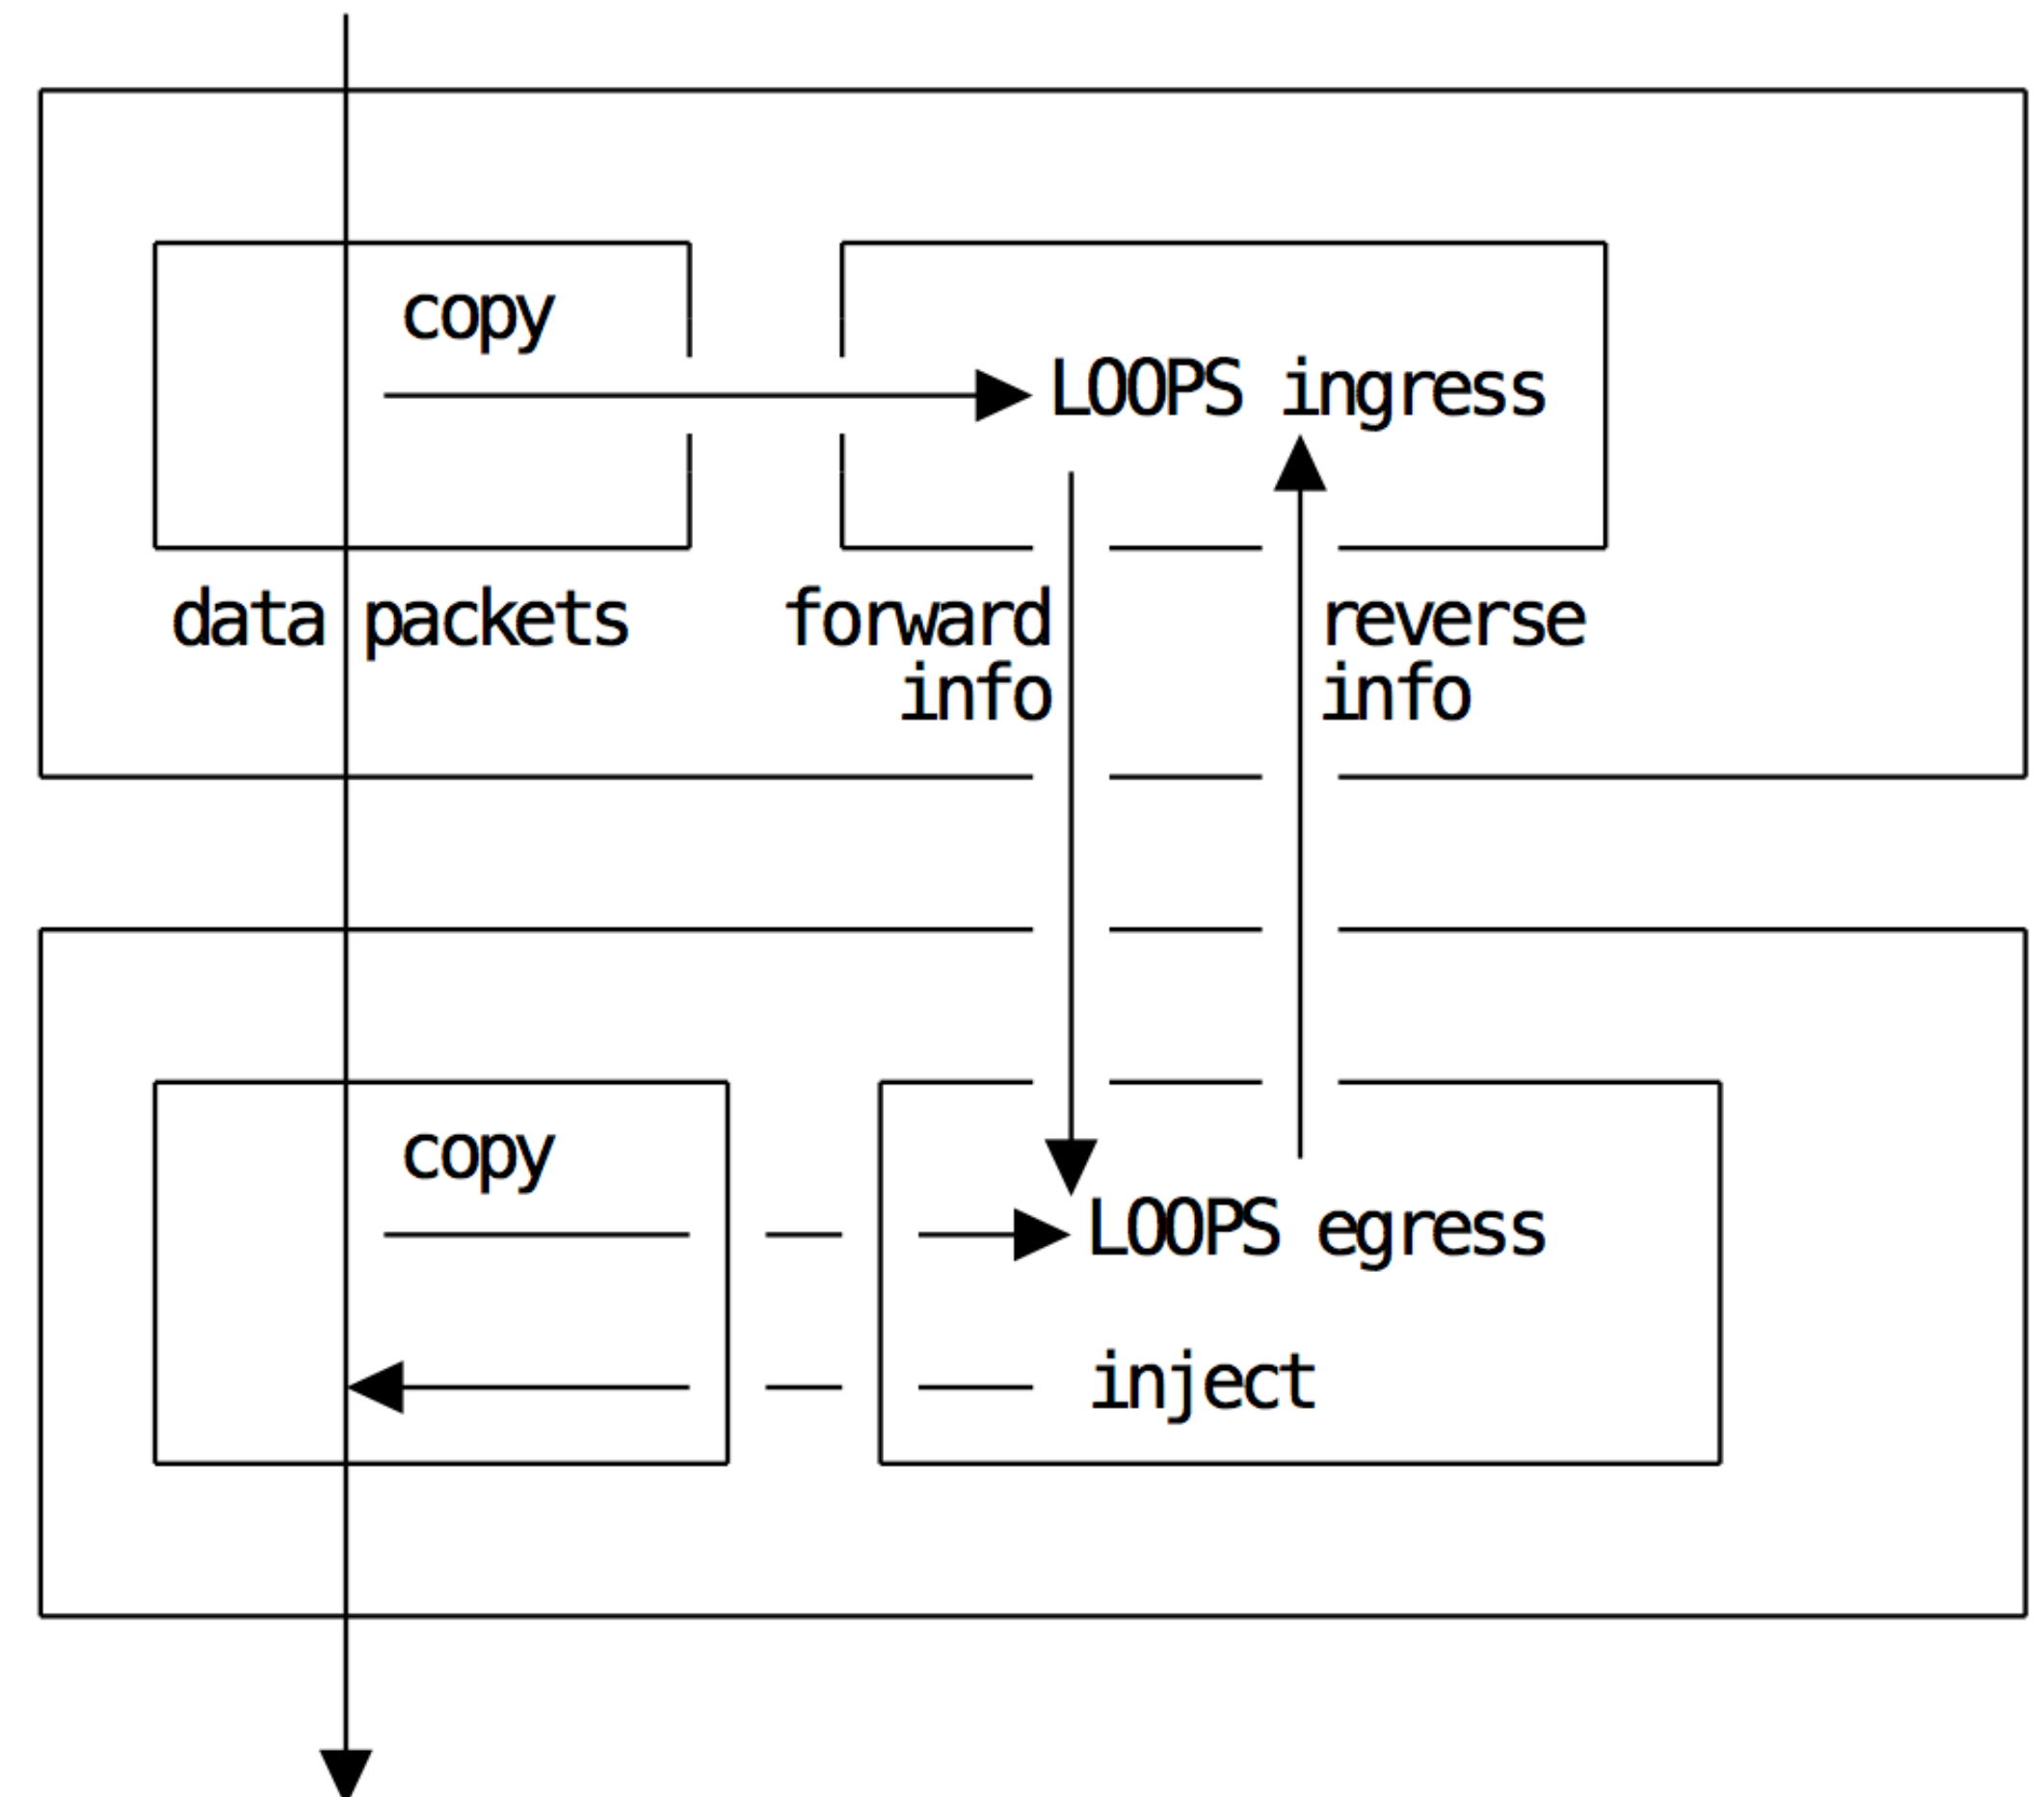


### 3. Ingress retransmits

- Detects need for retransmit  
**via timing only**
  - Decide based on congestion estimation as before
- Cost of hash collisions is low: misses retransmit opportunity.

### 4. Egress forwards

- That's all it does. There will be re-ordering.



# Summary: information exchanged

- **Forward:** nothing extra
- **Backward:** roughly as before - optional blocks type 1...  
*(can be piggybacked, aggregated, interspersed, repeated, ...)*
  - **Block 1**  
*(limited in some way: was optional, only upon "ACK desirable" for tunnel mode, but egress doesn't get this information in transparent mode)*
    - PSN being acknowledged
    - Absolute time of reception for the packet acked (**hash**)
  - **Block 2** (hmm)
    - an ACK Bloom filter?



# Conclusion

- Spectrum of possibilities, from "full-fledged" to min-intrusive
  - Various trade-offs between these options
- In all cases: LOOPS can be very beneficial when the LOOPS segment RTT is much shorter than the e2e RTT
  - Wireless NICs use this fact
- Some packet drops really cause pain
  - LOOPS can help




**Tail  
loss!**

# Recent work:

## Markov chain analysis of a LOOPS ingress cache

- System constraints: "transparent mode"
  - sender S emits packets with rate  $\lambda$
  - cache C additionally stores a copy of packets
  - loss detector L confirms reception of each packet
  - C retransmits if timer  $rto_c$  expires
  - R acknowledges packets to S
  - $p_i$ : packet loss probability in segment i




- Continuous-Time Markov Chain (CTMC) with finite states
  - Cache can be modelled as an M/D/1/N queuing model
  - Poisson arrival, deterministic service time
- Runa Barik, Michael Welzl, Peyman Teymoori, Safiqul Islam, Stein Gjessing: "Performance Evaluation of In-network Packet Retransmissions using Markov Chains", accepted for publication, CNC'20 workshop, proceedings of IEEE ICNC 2020, HI USA 2020.




# Results

- Major influencing factors:  $p_2$ ,  $rtt_c/rto_s$
- Non-linear dependence between cache filling level and packet loss probability  $p_2$
- Cache size has the least impact on its utilization: irrespective of the size, it can be fully utilized by higher packet loss probabilities
- LOOPS reduces retransmission rate and expected caching time at S

$p_1 = 0.01$ ,  $p_3 = 0.01$ ,  $rtt_s = 0.1$ ,  $rto_s = 0.2$ ,  $rtt_c = 0.05$ ,  $rto_c = 0.1$ , initial  $p_b = 0.01$



(a)  $\lambda = 5$



(b)  $\lambda = 10$

# Discussion items:

- Acks and retransmission:  
How does acknowledgement information + measurements flow to enable retransmission decisions (whether, when, how often)
- Encapsulation:  
Using encapsulation level formatting vs. our own (e.g. TLV, flag-based)  
Discuss Geneve as a starting point
- FEC framework:  
What is the benefit of new-fangled (e.g., RLC/sliding window) over traditional FEC (e.g., RaptorQ)

# Acks and retransmission

- LOOPS defines a protocol: What information do the two ends make available to each other when
- Retransmission decision can be local matter, if protocol supplies needed information
  - Traditional DUPACK-style: Any (small number of) packets that are ahead of the current order indicate a loss
  - Time-based: tolerate reordering; use more information than just

# Acks and retransmission

- Design so far:
  - PSN numbers payload packets (not LOOPS packets)
  - Block 1 ACKs (timely information): Latency & RTT measurement
    - **Ingress** decides whether this should be returned (ACK desirable bit)
    - Never requested for retransmissions (avoid ACK ambiguity)
  - Block 2 ACKs (aggregate information): ACK bitmaps
    - Basis for retransmission decisions
    - **Egress** decides when to acknowledge and how much, how often

# Acks and retransmission: What to achieve?

- Local retransmission most useful if it preempts end-to-end retransmission  
→ retransmission should happen quickly
  - But not so quickly that it creates spurious local retransmissions
  - Assumption: deadline for retransmission a couple RTTs after initial
  - One implementation strategy: Single retransmission only
- Potential goal: Robustness to some packet reordering on the path segment
- Class 1:
  - in average, significantly more than one packet is in flight
  - In average, loss recovery  $\ll 10\%$

# Acks and retransmission: Detect holes

- with reordering, hole-based retransmission schemes (e.g., Dupack) will cause some spurious retransmissions
  - end-to-end: too bad.
    - CC effect (reduction of CC window) will be canceled out quickly by spurious retransmission detection.
    - end-to-end delivery still faster with local recovery.
    - RACK etc. will ease this consideration over time.
  - LOOPS: Maybe don't do that, then?
- How does a time-based retransmission scheme look like for LOOPS?

# Time-based recovery (Strawman): RACK-inspired equivalent for LOOPS?

- Divide time into slots, e.g., milliseconds (or  $\sim$  an order of magnitude under min\_RTT; possibly defined in setup protocol)
- For each packet, remember slot when last sent
  - and whether already retransmitted (no longer used for measurement)
- Use aggregate ACKs to mark some slots as “mostly arrived”
- Subtract a reordering window, based on how hard that boundary is
- Pace out older unacknowledged packets at retransmission pace
- Flush slot when all acknowledged or beyond retransmission deadline

# Encapsulation

- LOOPS focuses on the “generic information” interchanged
- Need to map to specific encapsulation protocols.  
These have:
  - Required fields (that may map to LOOPS information)
  - Existing options (that may map to LOOPS information)
  - Extension points (TLVs, flag-based)

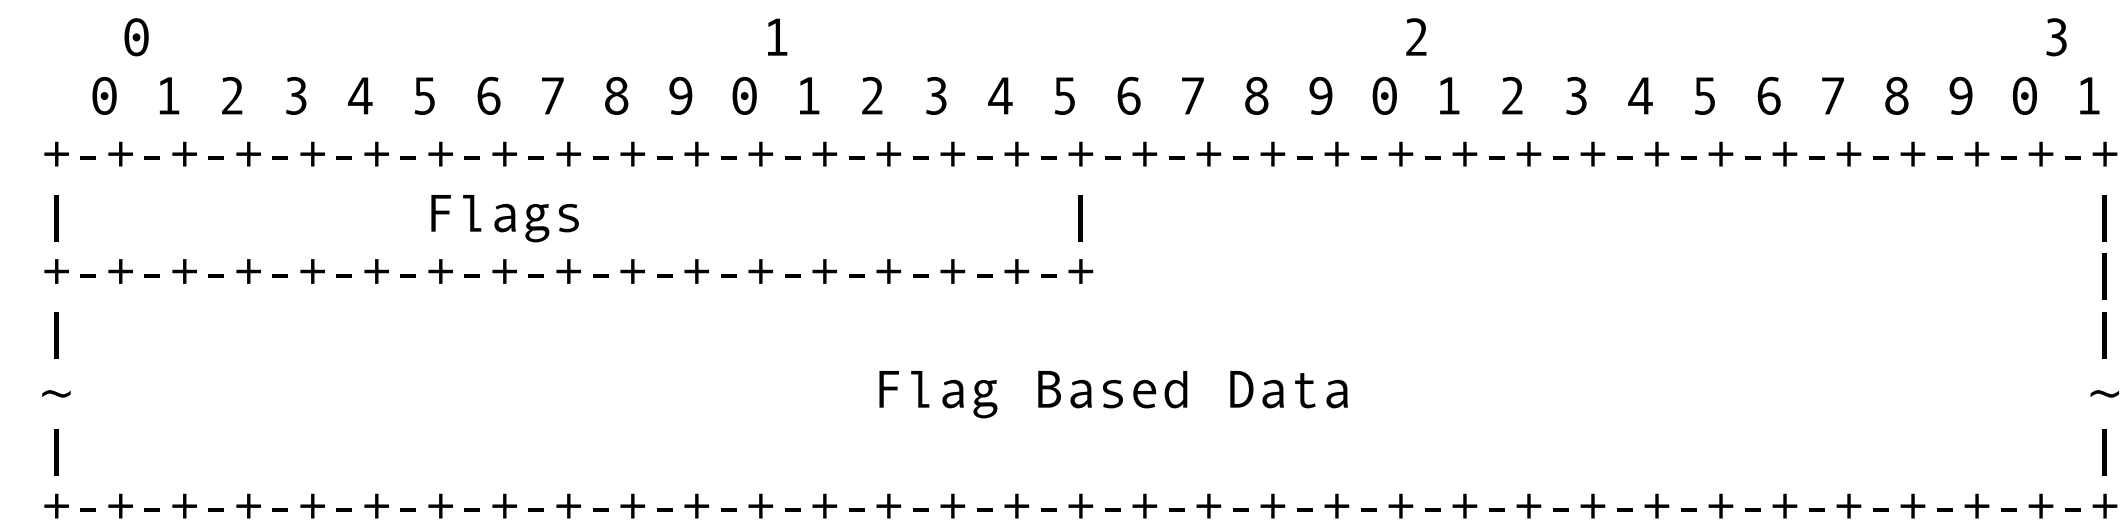


# Encapsulation

- LOOPS information may go into
  - Required fields, existing options (possibly after some mapping)
  - Extension points (via one-to-one mapping)
  - Blobs that go into an extension point (e.g., Geneve “LOOPS option”)
    - Need internal structure for representing those data:
      - LOOPS-specific approach, vs.
      - Mirror some of the encapsulation protocol’s approaches

# Encapsulation: Geneve Strawman


- draft-bormann-loops-geneve-binding-00.txt
- Goes for single LOOPS option:
  - Geneve:
    - 4-byte aligned, increments of 4 bytes
  - Flags for PSN/Ack, timestamp blob/echo, reception time, ...  
(and a few flag-only flags, e.g., ACK desirable)
  - Separate flag for (single) non-4-byte element: Block2
- Unfortunately, the items themselves then are not 4-byte aligned



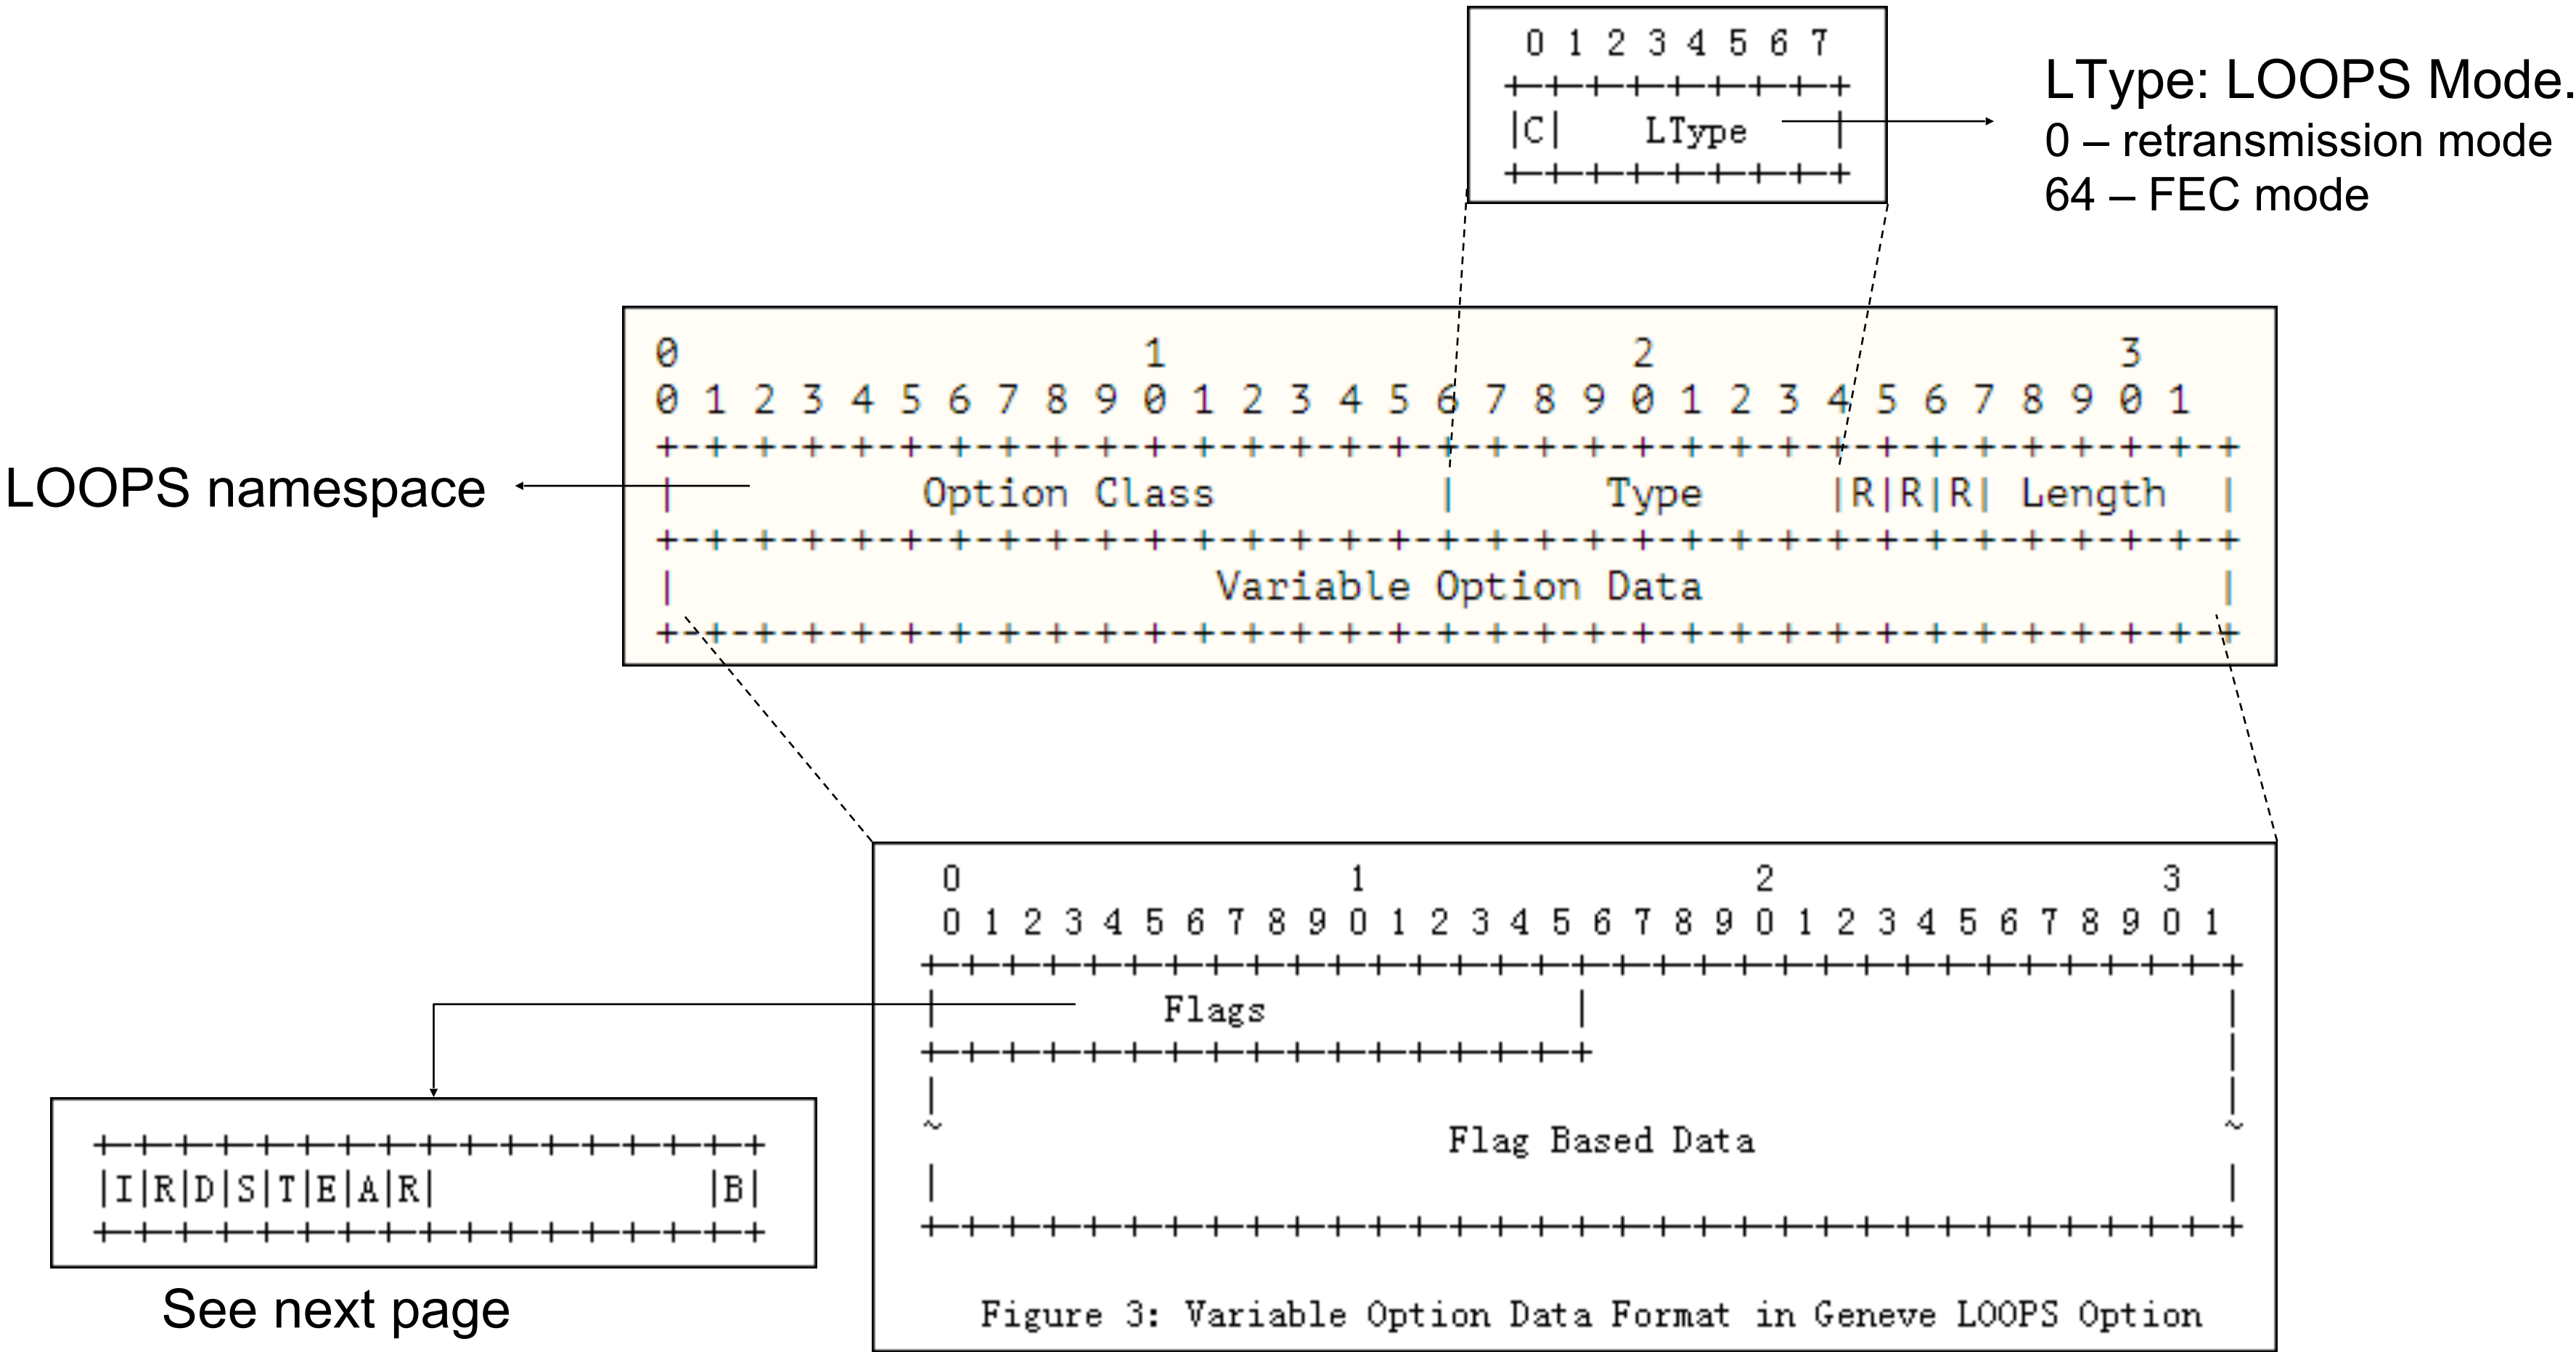
# Geneve Overview


- Tunnel protocol primarily used in DC to support virtualization
- UDP based: X over UDP over IP
- Data plane extensibility: by TLVs
- Pure data plane spec. Can leverage many control planes.
- Two categories
  - Tunnel endpoint
    - Originate/terminate packets
    - Can insert/delete/parse options
    - **LOOPS should work on these points**
  - Transit device
    - Can interpret options
    - Must not insert/delete/change options
    - Must not modify Geneve headers

# Geneve Format



# Proposed LOOPS option - map LOOPS to tunnel protocol





# Flags in LOOPS option

Used on their own

- I: Initial Packet Sequence Number (PSN) flag; may be set by the LOOPS ingress to notify the egress about using a new initial PSN.
- R: Initial PSN Received flag; echo of I flag provided by the LOOPS egress.
- D: ACK Desired flag; set by the LOOPS ingress if it wants the egress to generate an acknowledgement immediately upon receiving a particular packet.

Used with additional 32-bit in “flag based data” field

- S: PSN flag; indicates a PSN data block is carried. It must be set when a packet payload is present. It must not be set if the packet is a pure LOOPS ACK packet, i.e. when no payload is included in the packet.
- T: Timestamp flag; indicates a Timestamp data block is carried
- E: Echoed Timestamp flag.
- A: ACK number flag; indicates presence of a BLOCK 1 ACK info
- R: Reception time flag; May only be set if A is set. Indicates that an absolute reception time is carried.

Used with additional variable-length block in “flag based data” field (least bit)

- B: Block 2 ACK info flag; indicates presence of a BLOCK 2 ACK info

# Design choices and open questions

1. Single option or Multiple options
  - Single: Efficient, compact
2. Pure ACK
  - Set O bit in Geneve header
  - A control packet without payload
3. VNI:
  - mandatory in Geneve
  - A specific VNI can be used for LOOPS enabled traffic, or for a particular tenant's LOOPS enabled traffic, or
4. Alignment: Geneve requires 4 byte alignment. PSN length?
5. Move some flags to “Type” field of generic option header?
6. Length of Timestamp & Echoed Timestamp?

# FEC approach

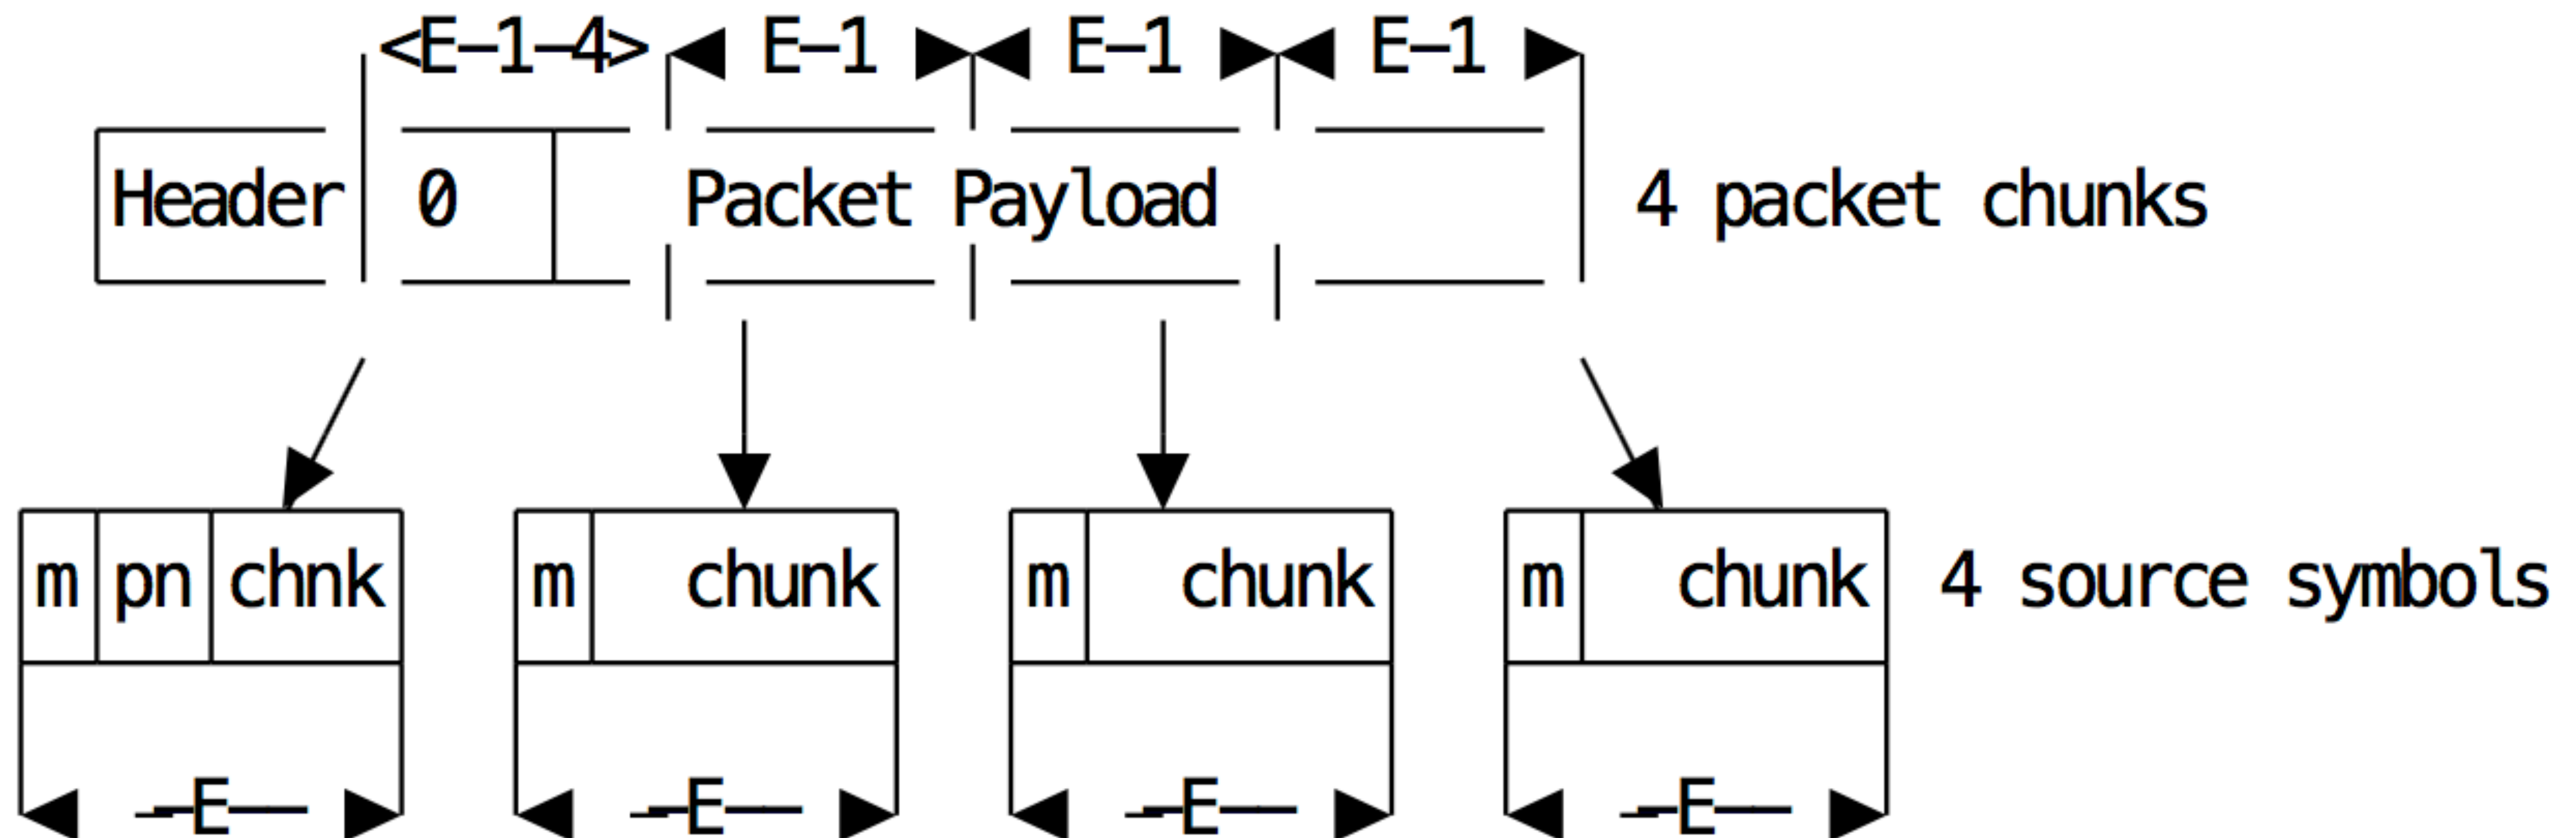
- Support multiple classes of FEC schemes, e.g.:
  - Very simple parity (as in SMPTE 2022)
  - Fountain Codes (e.g., RaptorQ)
  - Sliding Window schemes (e.g., RLC)
- Assume all codes are systematic (needed for transparent mode)
  - Except for transparent mode, augment payload packets by FEC indices
- Possibly add special handling for larger-than-tunnel-MTU packets
- Add repair packets with repair information



# FEC approach

- LOOPS can provide:
  - forward: place for FEC indices, separate format for repair packets
  - return: Block 2 acknowledgements, or aggregate loss rate feedback
- Assumption: large size variance of payload packets (avg 400..700 B)
  - Payload packets are divided up before being funneled into FEC
  - Not necessarily related to the way they are sent forward
  - Any piggybacking for repair segments?  
Recombining/splitting of payload packets (also for MTU reasons)?

From draft-roca-nwcrg-rlc-fec-scheme-for-quic-02:



# FEC: Design choices

- Classes of FEC schemes (that can be handled equivalently by LOOPS)
  - What are the FEC indices to be added to payload packets?  
(Tunnel: right there; Transparent: separately)
- Do we put in some MTU mitigation (breaking up payload packets)?  
Piggy-backing runts/short packets/repair symbols?
- Feedback:
  - For controlling FEC rate — what is the time scale?
  - For filling in repair packets?
- Details of the construction of FEC input and repair packets;  
how are reconstructed packets put together again?

# Next Steps

# While we are not a WG...

- Continue on, *working* like a WG
  - Explore design space, maybe holding back on tough decisions for now
- Continue improving the set of documents, possibly adding FEC document
  - Identify authors and reviewers
- Employ [github.com/loops-wg](https://github.com/loops-wg) and [loops@ietf.org](mailto:loops@ietf.org) for coordination
- Review charter proposal at github; react to AD input on this
- Aim for being a WG at IETF 107 (Vancouver, March 2020)